

Sockets and JDBC

Weng Kai

Identifying a machine

- * 两种形式的地址：
 - The familiar DNS (Domain Name Service) form like `www.zju.edu.cn`.
 - The "dotted quad" form, which is four numbers separated by dots, such as `123.255.28.120`.
- * Case Study:WhoAml.java

TCP vs. UDP

* TCP是有连接的协议

- 在数据传送前要先建立端到端的连接，在数据传送过程中会校验数据

* UDP是无连接的

- 数据传送前不需要建立连接，不管对方是否存在可以盲发

Port


NAT

socket port

套接字, 64K 监听


PDP-11

16-bit
big-endian



server

RDD



TCP

- * 是一种Client/Server模式

- * 端口

- 每个IP设备最多可以有65536个端口，通信是在端口与端口之间进行的

- * Server方在固定的端口守听，client方连接该端口，从而形成数据链路

TCP Sockets

* ServerSocket

- 守听在固定端口，等待client连接的对象

* Socket

1. Client方用来连接Server方的对象
2. Server方用来和Client连接的对象

数据传输

* 在Socket对象里有InputStream和OutputStream, 用来传输数据

* nc -l 12345 —> listen

* nc 12345 —> connect

* Case Study:

* 单服务

- JabberServer.java

- JabberClient.java

* 多服务

- MutiJabberServer.java

- MutiJabberClient.java

UDP

* DatagramSocket

- UDP端口，收或发

* DatagramPacket

- UDP数据


* Case:

- DayBcast.java
- DayWatch.java

RTTI in Socket Server

- Binary protocol usually use the first byte as type of messages
- Sub-classes should be used to deal with different types of messages
- Build a Hash-Map of sub-classes Class object to generate dealer of messages

JDBC



* Java DataBase Connection

* 遵循ODBC模型，但不是ODBC

JDBC三部曲

- * DriverManager

- * Connection

- * Statement

- * ResultSet

- * Case Study: Lookup.java

更新数据库

* `Statement.executeUpdate(String SQL);`

PreparedStatement

- * 由Connection产生PreparedStatement对象，其中的参数用?表示
- * 用PreparedStatement的setXXX()函数对参数赋值
- * 用PreparedStatement的execute()执行
- * Case: Query DB.java

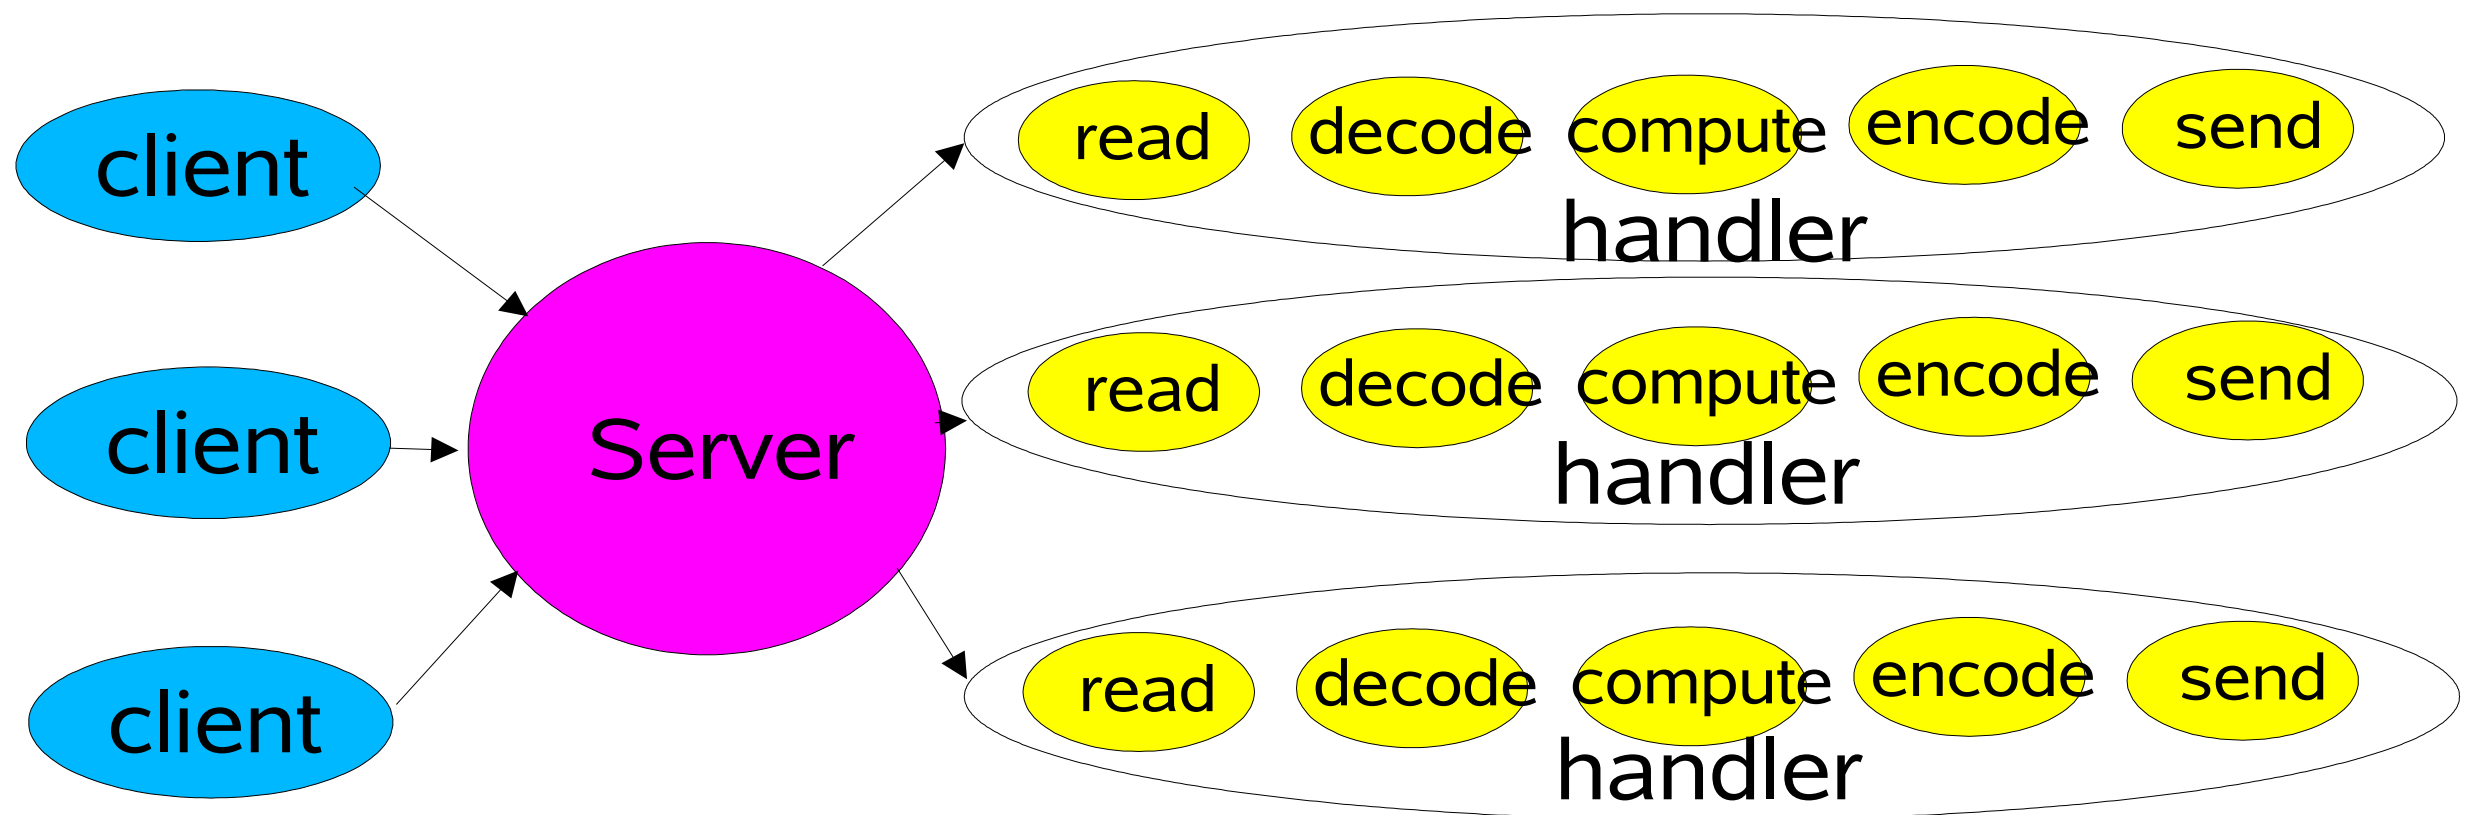
事务处理

- * 在Connection上做事务处理
- * `setAutoCommit()`
- * `commit()`
- * `rollback()`

Network Services

- Most have same basic structure:
 - Read request
 - Decode request
 - Process service
 - Encode reply
 - Send reply
- But differ in nature and cost of each step
 - XML parsing, File transfer, Web page generation, computational services, ...

Classic Service Designs



Classic ServerSocket Loop

```
class Server implements Runnable {  
    public void run() {  
        try {  
            ServerSocket ss = new ServerSocket(PORT);  
            while (!Thread.interrupted())  
                new Thread(new Handler(ss.accept())).start();  
            // or, single-threaded, or a thread pool  
        } catch (IOException ex) { /* ... */ }  
    }  
}
```

```
static class Handler implements Runnable {  
    final Socket socket;  
    Handler(Socket s) { socket = s; }  
    public void run() {  
        try {  
            byte[] input = new byte[MAX_INPUT];  
            socket.getInputStream().read(input);  
            byte[] output = process(input);  
            socket.getOutputStream().write(output);  
        } catch (IOException ex) { /* ... */ }  
    }  
    private byte[] process(byte[] cmd) { /* ... */ }  
}
```

多机聊天时的发送问题


- 一个client发送上来的一条消息需要发送给所有的其他client（通常也包括这个client）
- 每个client有一个发送线程（在接收线程之外）和一个发送队列，有一个中央机制遍历全部发送队列送消息
- 有一个中央的发送线程和队列，遍历全部client发送消息
 - 可能被某个client堵住
 - 为每个client起一个线程发送一条消息

Scalability Goals

- Graceful degradation under increasing load (more clients)
- Continuous improvement with increasing resources (CPU, memory, disk, bandwidth)
- Also meet availability and performance goals
 - Short latencies
 - Meeting peak demand
 - Tunable quality of service
- Divide-and-conquer is usually the best approach for achieving any scalability goal

Divide and Conquer

- Divide processing into small tasks
 - Each task performs an action without blocking
- Execute each task when it is enabled
 - Here, an IO event usually serves as trigger

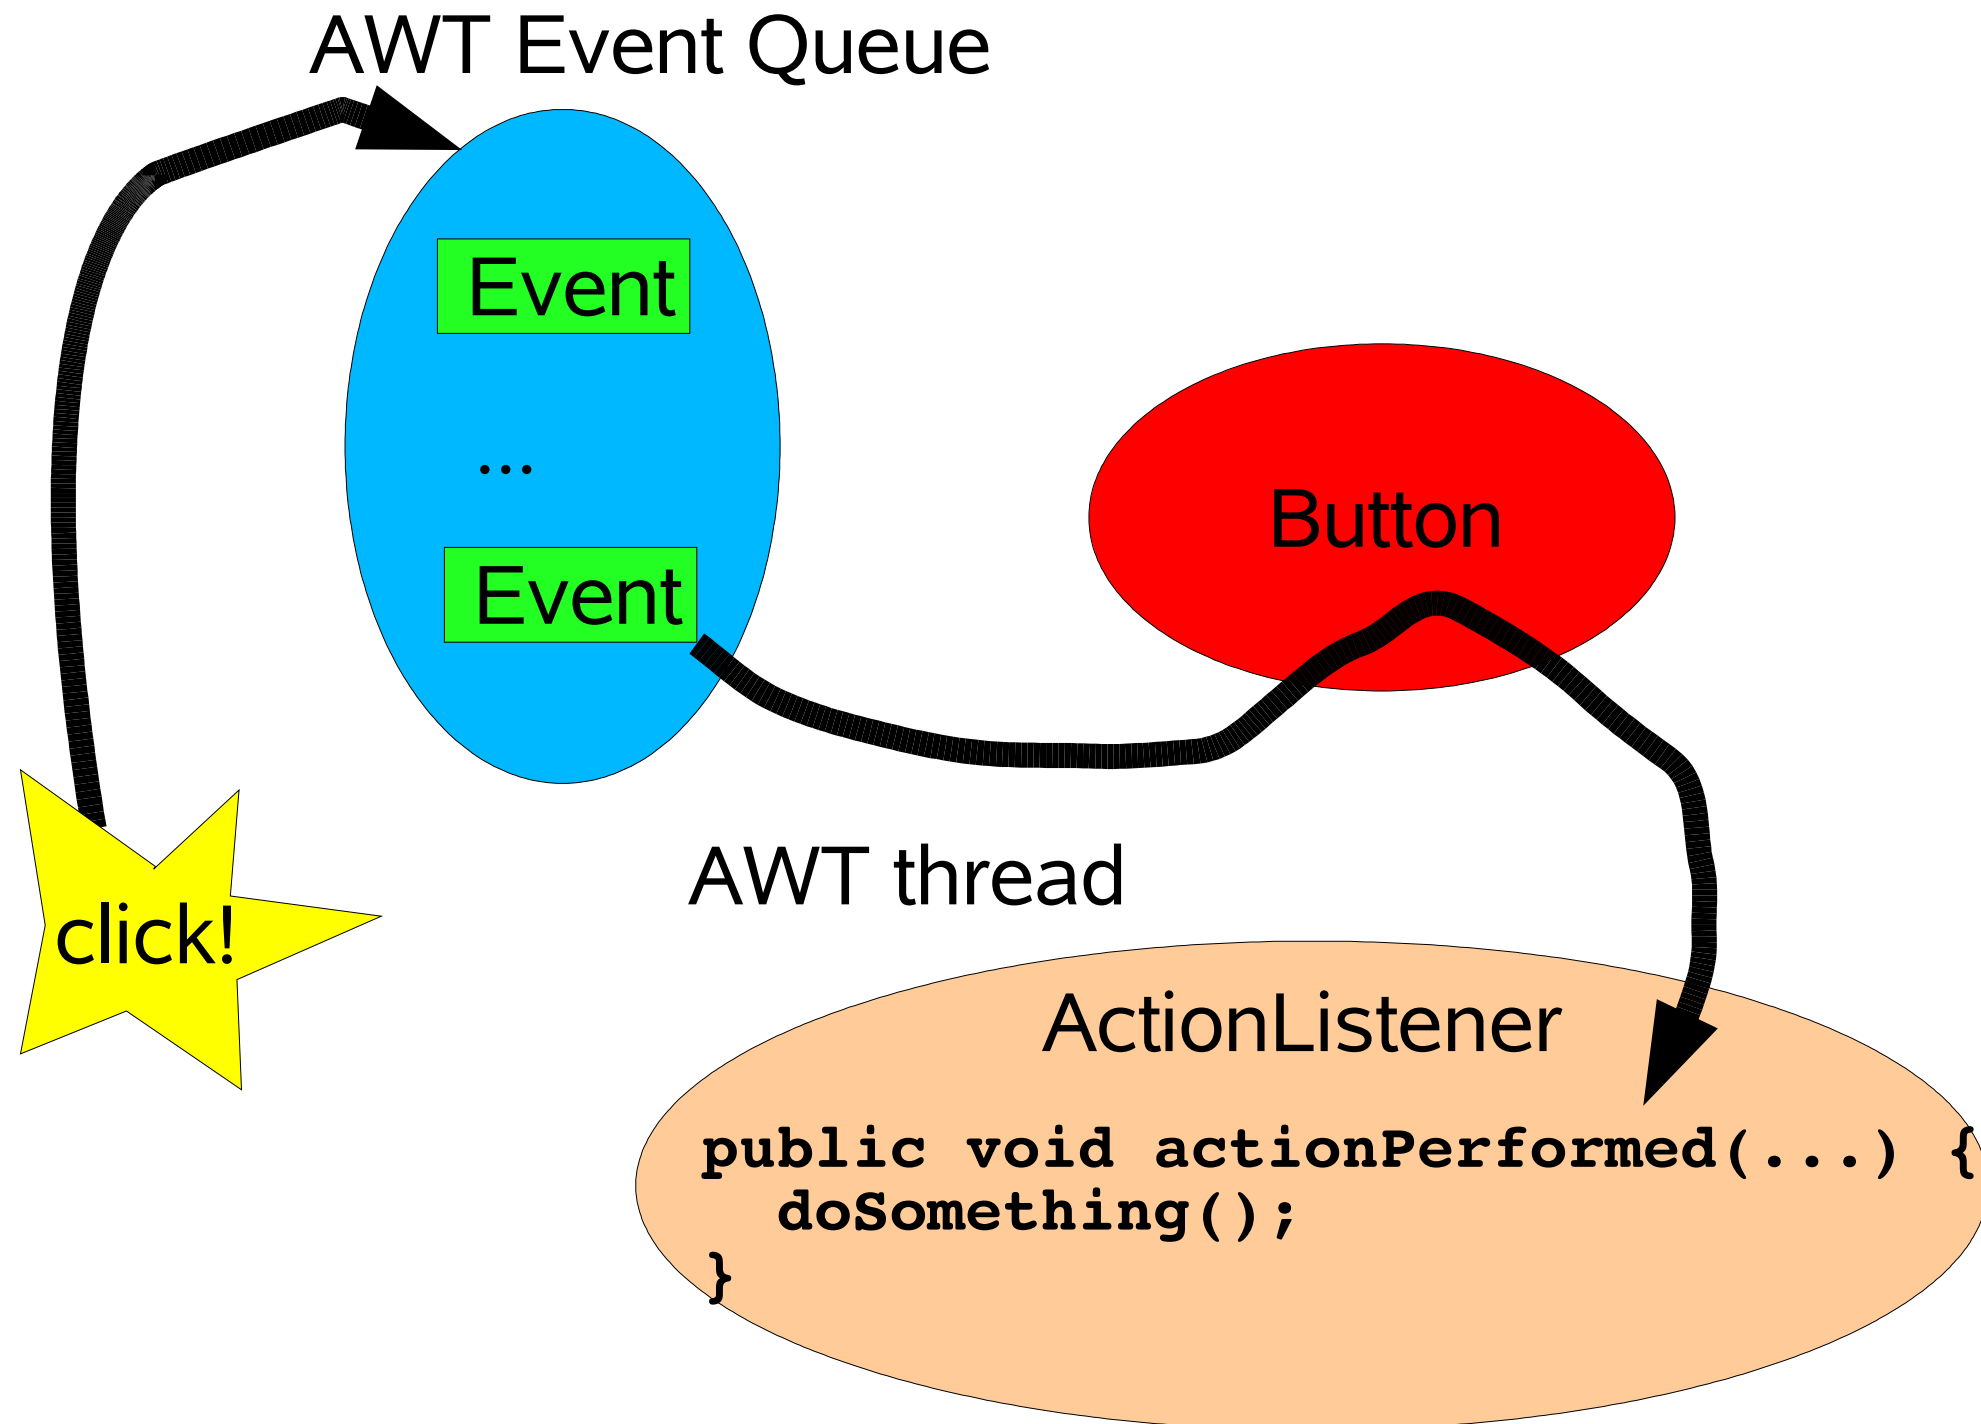


- Basic mechanisms supported in java.nio
 - Non-blocking reads and writes
 - Dispatch tasks associated with sensed IO events
- Endless variation possible
 - A family of event-driven designs

Event-driven Designs

- Usually more efficient than alternatives
 - Fewer resources: Don't usually need a thread per client
 - Less overhead: Less context switching, often less locking
 - But dispatching can be slower: Must manually bind actions to events
- Usually harder to program
 - Must break up into simple non-blocking actions
 - Similar to GUI event-driven actions
 - Cannot eliminate all blocking: GC, page faults, etc
 - Must keep track of logical state of service


Background: Events in AWT



Reactor Pattern

- Reactor responds to IO events by dispatching the appropriate handler
 - Similar to AWT thread
- Handlers perform non-blocking actions
 - Similar to AWT ActionListener
- Manage by binding handlers to events
 - Similar to AWT addActionListener
- See Schmidt et al, Pattern-Oriented Software Architecture, Volume 2 (POSA2)
 - Also Richard Stevens's networking books, Matt Welsh's SEDA framework, etc

Basic Reactor Design



- Single threaded version

java.nio Support

- Channels: Connections to files, sockets etc that support non-blocking reads
- Buffers: Array-like objects that can be directly read or written by Channels
- Selectors: Tell which of a set of Channels have IO events
- SelectionKeys: Maintain IO event status and bindings

Reactor 1: Setup

```
class Reactor implements Runnable {
    final Selector selector;
    final ServerSocketChannel serverSocket;

    Reactor(int port) throws IOException {
        selector = Selector.open();
        serverSocket = ServerSocketChannel.open();
        serverSocket.socket().bind(
            new InetSocketAddress(port));
        serverSocket.configureBlocking(false);
        SelectionKey sk =
            serverSocket.register(selector,
                SelectionKey.OP_ACCEPT);
        sk.attach(new Acceptor());
    }

    /*
    Alternatively, use explicit SPI provider:
    SelectorProvider p = SelectorProvider.provider();
    selector = p.openSelector();
    serverSocket = p.openServerSocketChannel();
    */
}
```

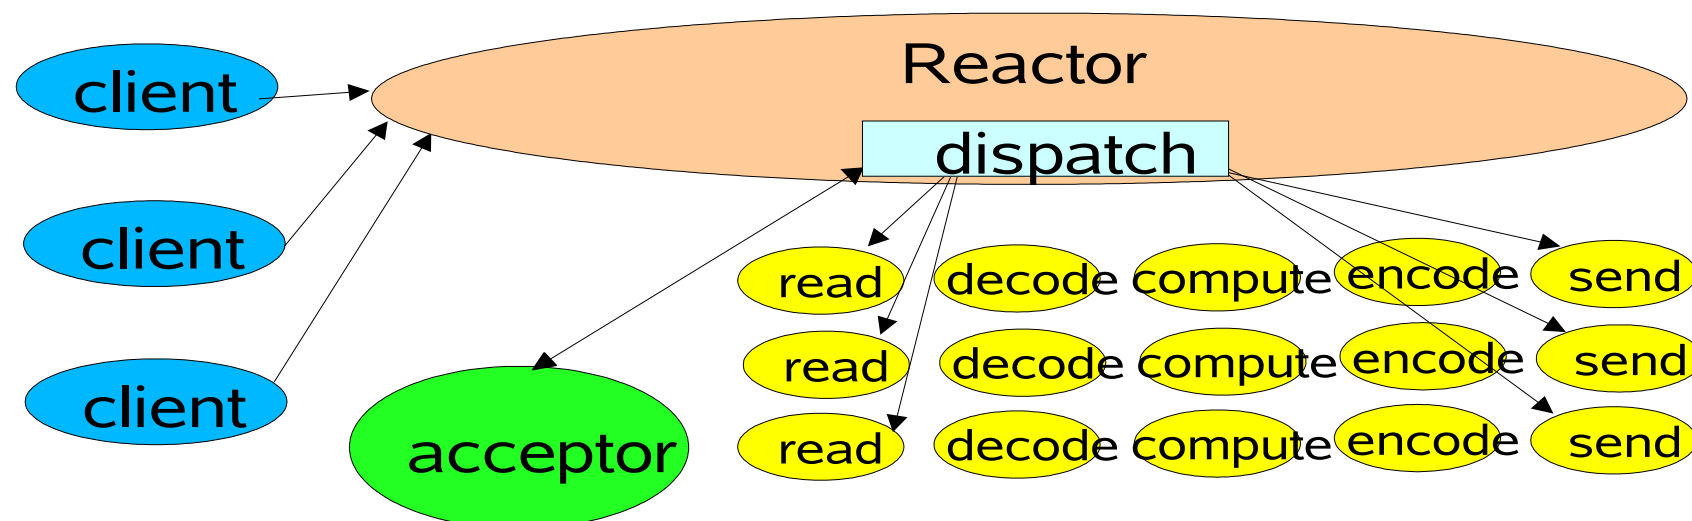
Reactor 2: Dispatch Loop

```
// class Reactor continued
public void run() { // normally in a new
Thread
    try {
        while (!Thread.interrupted()) {
            selector.select();
            Set selected = selector.selectedKeys();
            Iterator it = selected.iterator();
            while (it.hasNext())
                dispatch((SelectionKey) it.next());
            selected.clear();
        }
    } catch (IOException ex) { /* ... */ }
}

void dispatch(SelectionKey k) {
    Runnable r = (Runnable) (k.attachment());
    if (r != null)
        r.run();
}
```

Reactor 3: Acceptor

```
// class Reactor continued
class Acceptor implements Runnable { // inner
    public void run() {
        try {
            SocketChannel c = serverSocket.accept();
            if (c != null)
                new Handler(selector, c);
        }
        catch(IOException ex) { /* ... */ }
    }
}
```



Reactor 4: Handler setup

```
final class Handler implements Runnable {  
    final SocketChannel socket;  
    final SelectionKey sk;  
    ByteBuffer input = ByteBuffer.allocate(MAXIN);  
    ByteBuffer output = ByteBuffer.allocate(MAXOUT);  
    static final int READING = 0, SENDING = 1;  
    int state = READING;
```

```
    Handler(Selector sel, SocketChannel c)  
        throws IOException {  
        socket = c; c.configureBlocking(false);  
        // Optionally try first read now  
        sk = socket.register(sel, 0);  
        sk.attach(this);  
        sk.interestOps(SelectionKey.OP_READ);  
        sel.wakeup();  
    }
```

```
    boolean inputIsComplete() { /* ... */ }  
    boolean outputIsComplete() { /* ... */ }  
    void process() { /* ... */ }
```

Reactor 5: Request handling

```
// class Handler continued
public void run() {
    try {
        if (state == READING) read();
        else if (state == SENDING) send();
    } catch (IOException ex) { /* ... */ }
}

void read() throws IOException {
    socket.read(input);
    if (inputIsComplete()) {
        process();
        state = SENDING;
        // Normally also do first write now
        sk.interestOps(SelectionKey.OP_WRITE);
    }
}

void send() throws IOException {
    socket.write(output);
    if (outputIsComplete()) sk.cancel();
}
}
```

Per-State Handlers

```
class Handler { // ...

    public void run() { // initial state is reader
        socket.read(input);
        if (inputIsComplete()) {
            process();
            sk.attach(new Sender());
            sk.interest(SelectionKey.OP_WRITE);
            sk.selector().wakeup();
        }
    }

    class Sender implements Runnable {
        public void run(){ // ...
            socket.write(output);
            if (outputIsComplete()) sk.cancel();
        }
    }
}
```

- A simple use of GoF State-Object pattern
 - Rebind appropriate handler as attachment

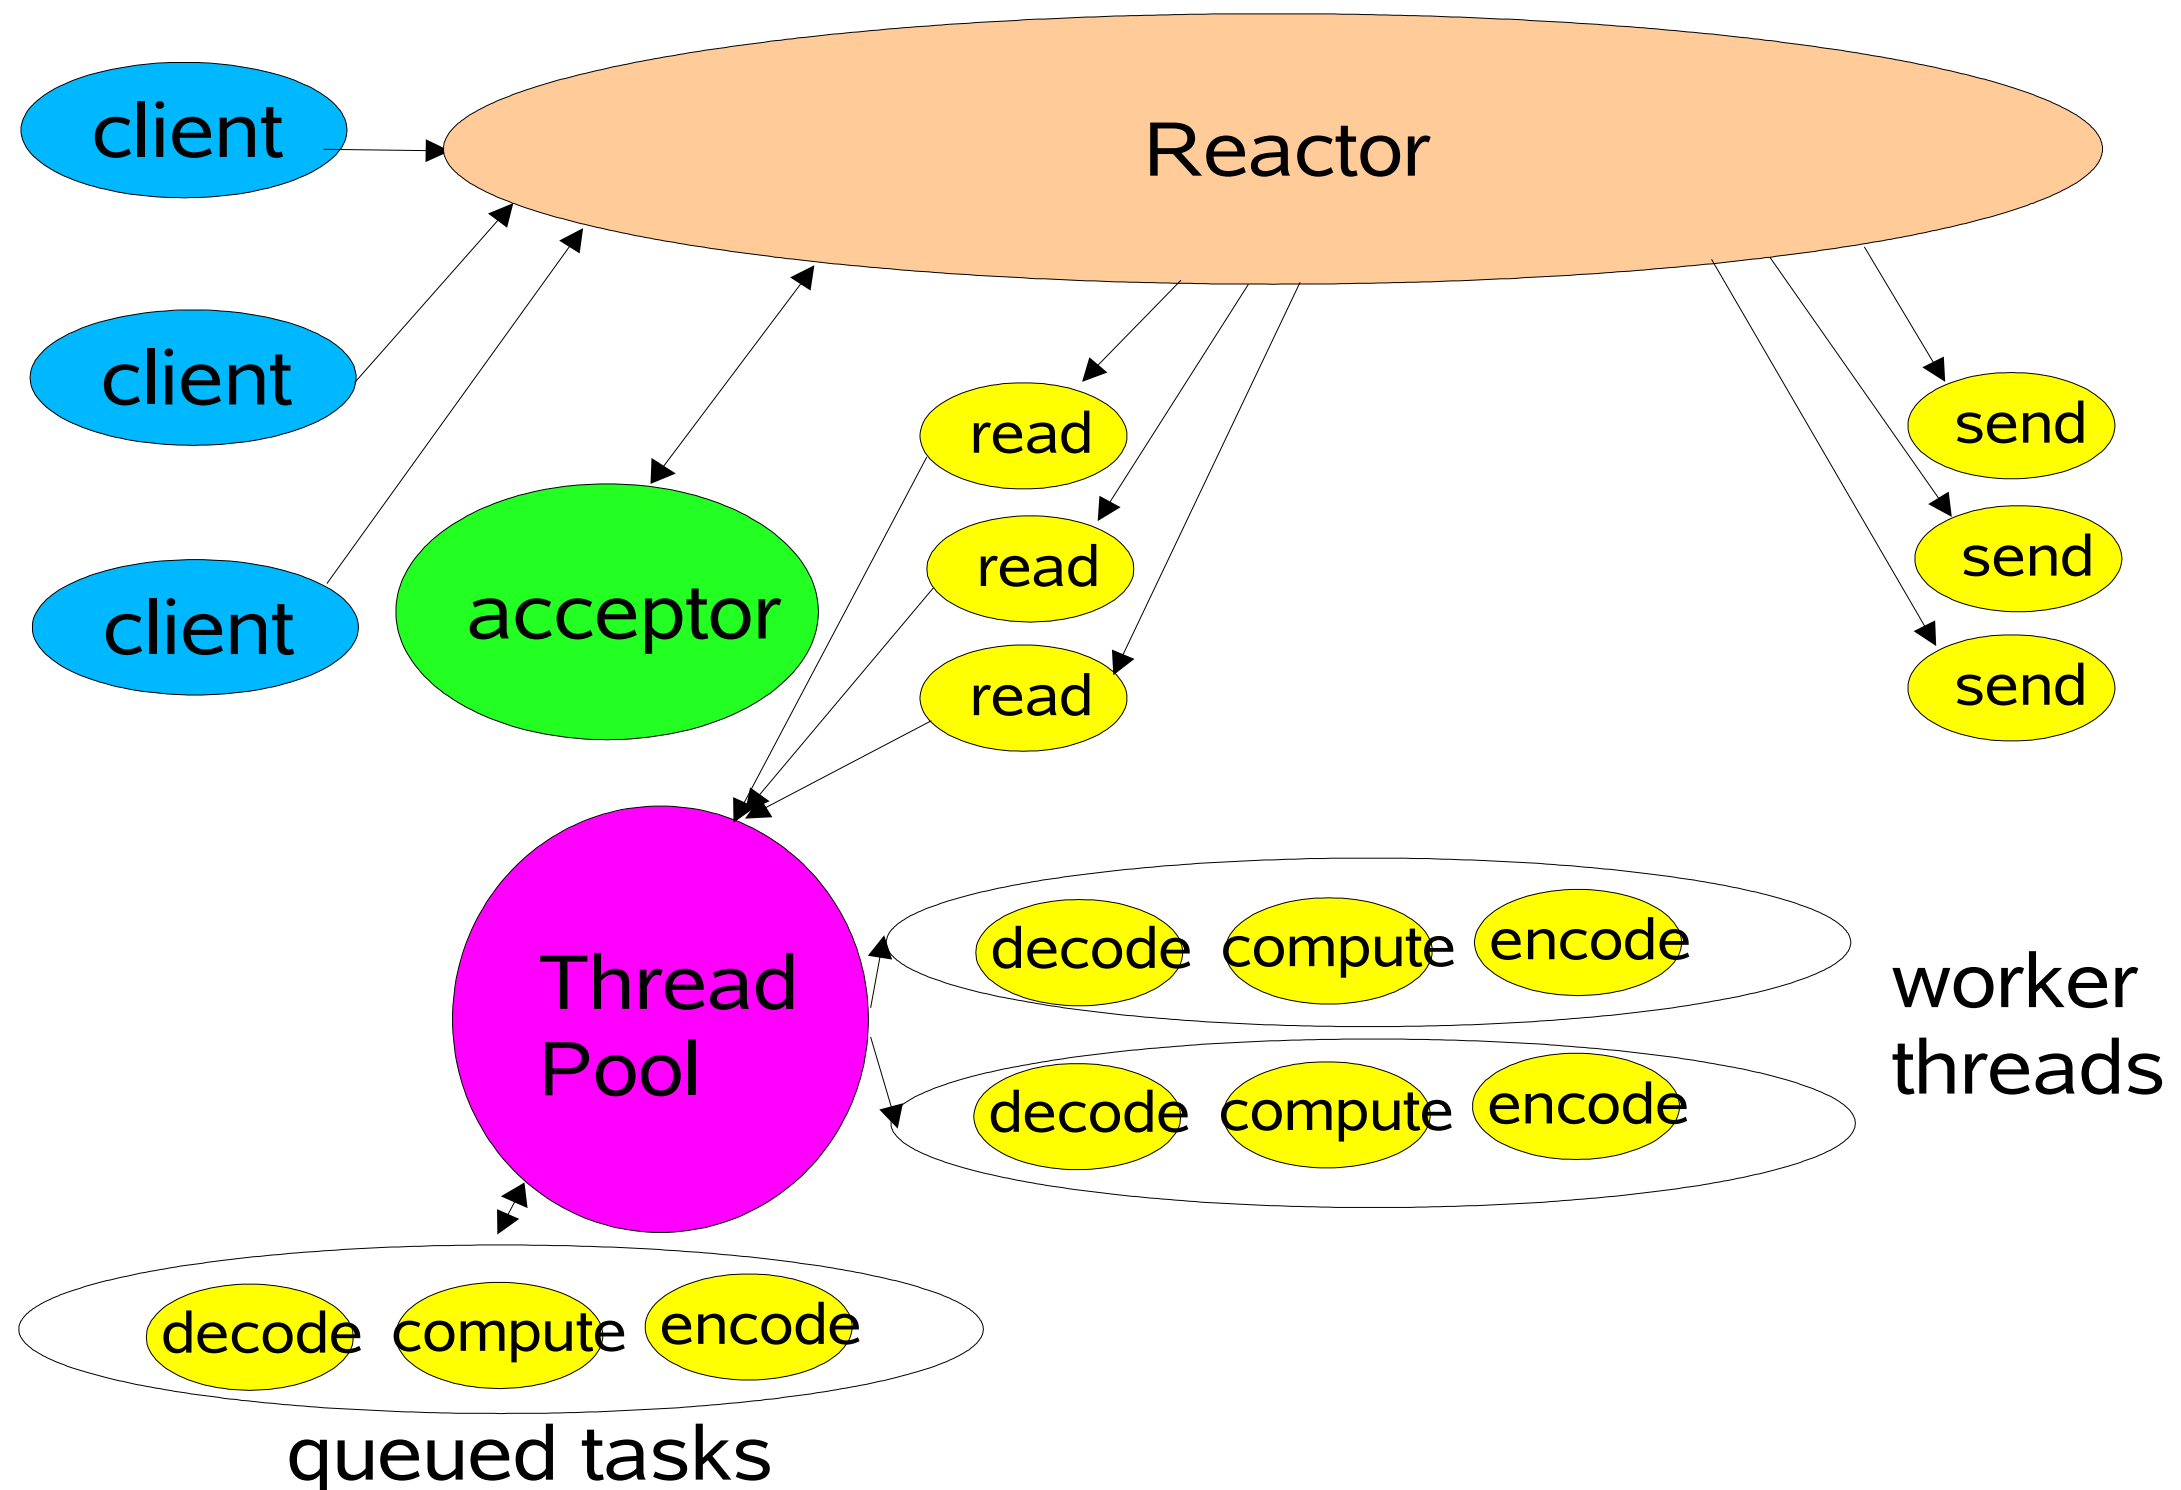
Multithreaded Designs

- Strategically add threads for scalability: Mainly applicable to multiprocessors
- Worker Threads
 - Reactors should quickly trigger handlers
 - Handler processing slows down Reactor
 - Offload non-IO processing to other threads
- Multiple Reactor Threads
 - Reactor threads can saturate doing IO
 - Distribute load to other reactors
 - Load-balance to match CPU and IO rates

Worker Threads

- Offload non-IO processing to speed up Reactor thread
 - Similar to POA2 Proactor designs
- Simpler than reworking compute-bound processing into event-driven form
 - Should still be pure nonblocking computation
 - Enough processing to outweigh overhead
- But harder to overlap processing with IO
 - Best when can first read all input into a buffer
- Use thread pool so can tune and control
 - Normally need many fewer threads than clients

Worker Thread Pools



Handler with Thread Pool

```
class Handler implements Runnable {
    // uses util.concurrent thread pool
    static PooledExecutor pool = new PooledExecutor(...);
    static final int PROCESSING = 3;
    // ...

    synchronized void read() { // ...
        socket.read(input);
        if (inputIsComplete()) {
            state = PROCESSING;
            pool.execute(new Processer());
        }
    }

    synchronized void processAndHandOff() {
        process();
        state = SENDING; // or rebind attachment
        sk.interest(SelectionKey.OP_WRITE);
    }

    class Processer implements Runnable {
        public void run() { processAndHandOff(); }
    }
}
```

Coordinating Tasks

- Handoffs
 - Each task enables, triggers, or calls next one
 - Usually fastest but can be brittle
- Callbacks to per-handler dispatcher
 - Sets state, attachment, etc
 - A variant of GoF Mediator pattern
- Queues: For example, passing buffers across stages
- Futures
 - When each task produces a result
 - Coordination layered on top of join or wait/notify

Using PooledExecutor

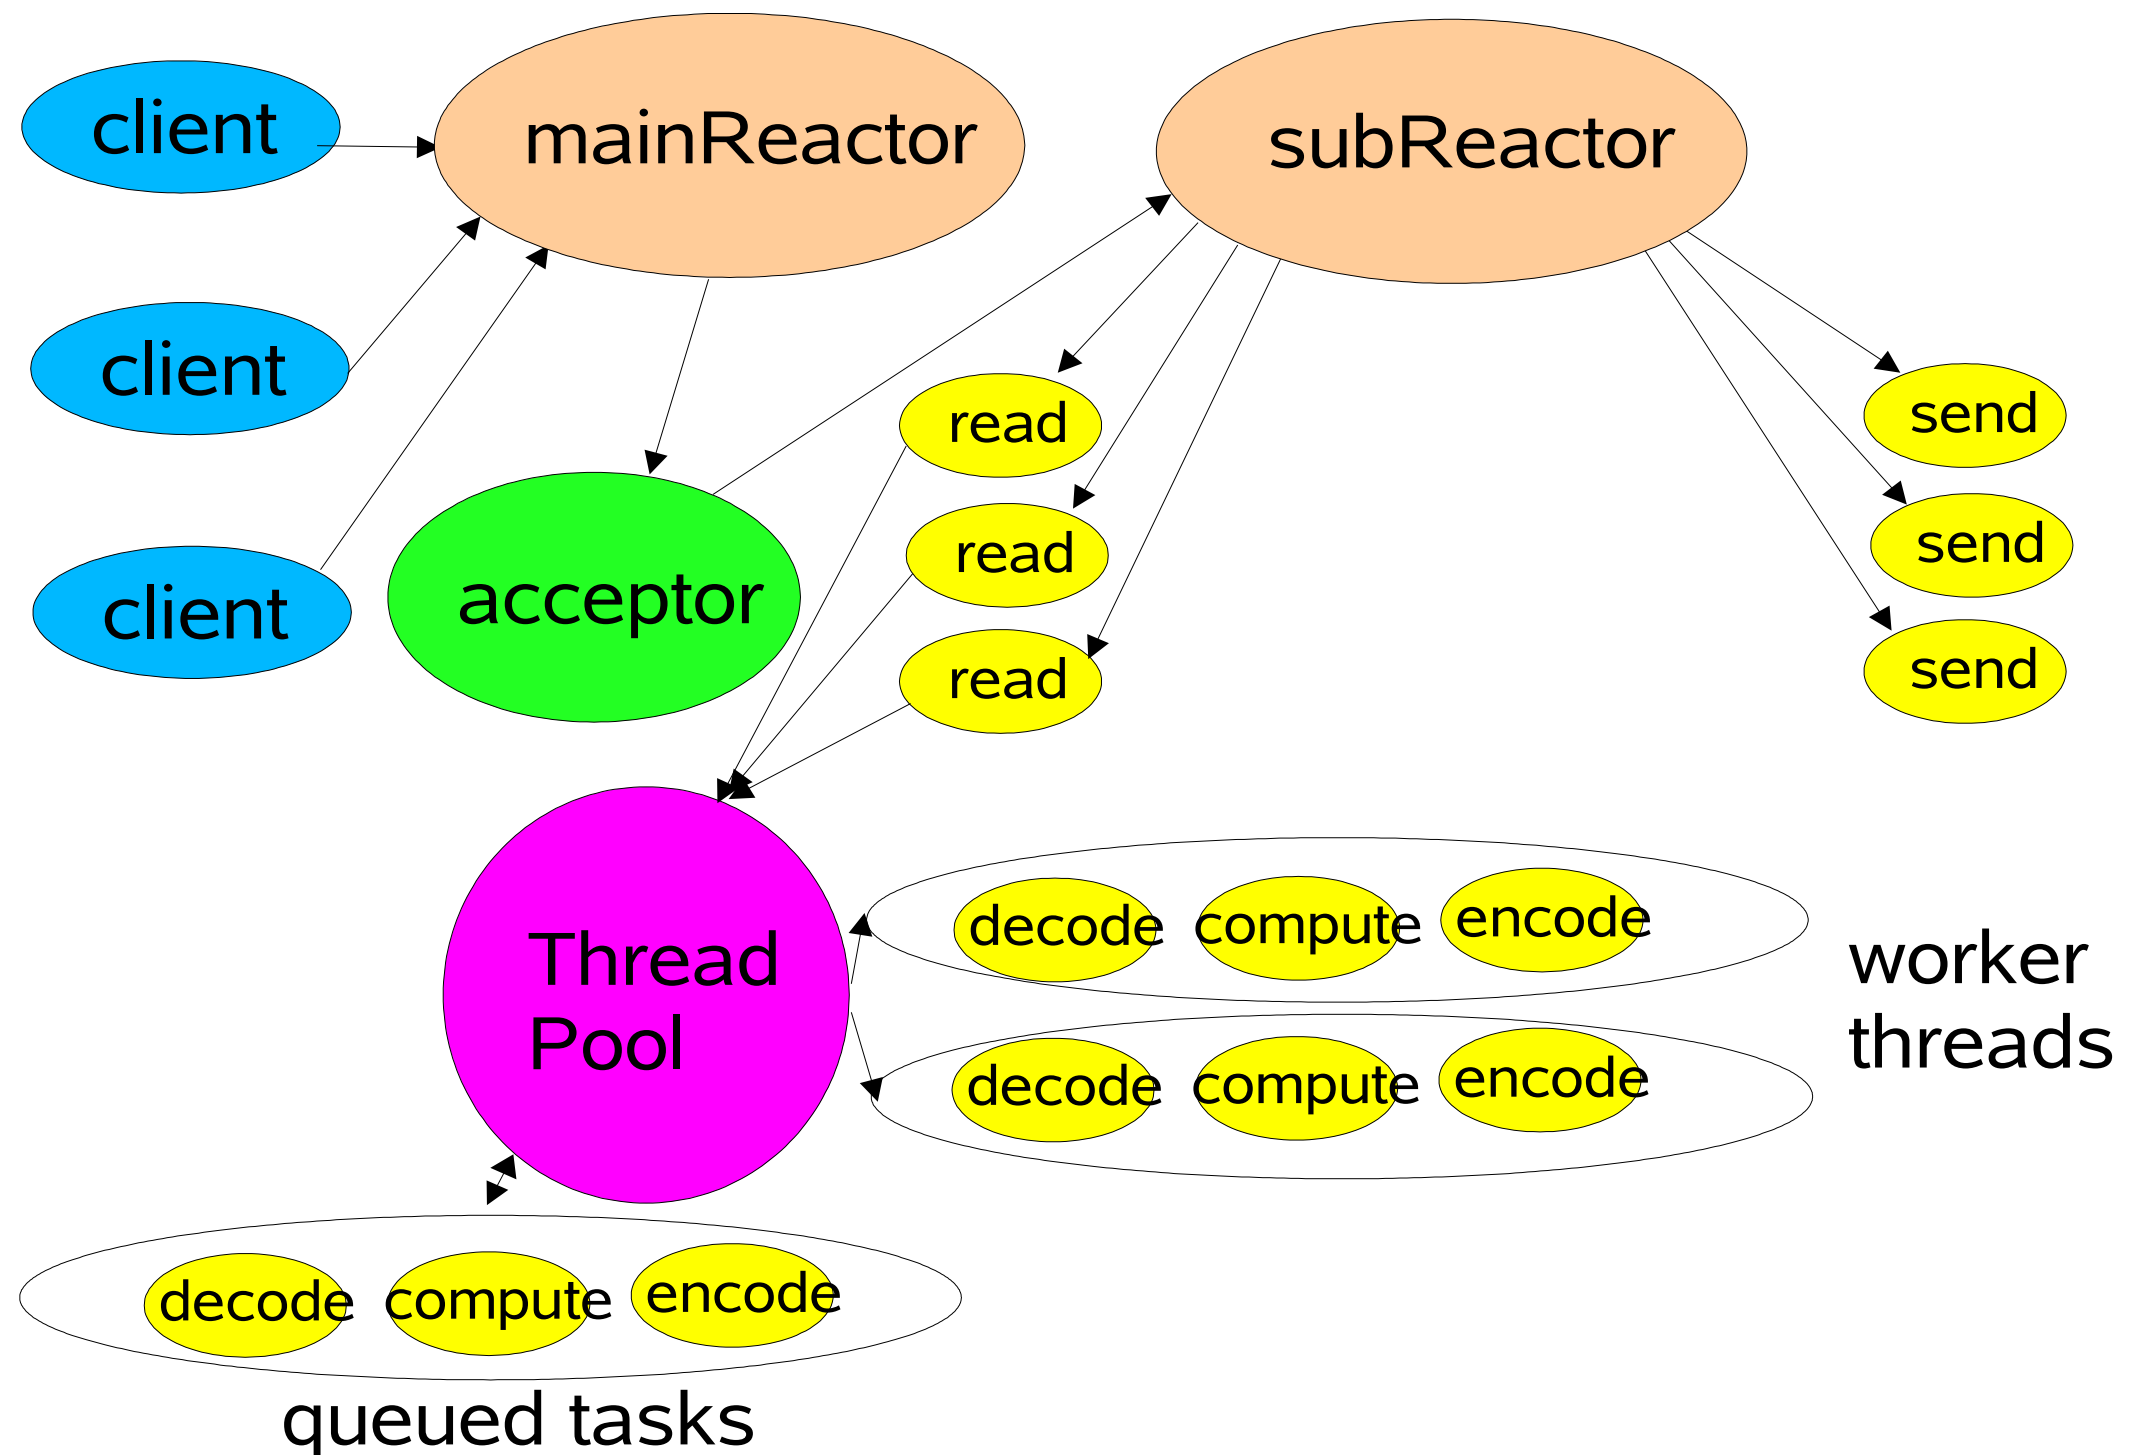
- A tunable worker thread pool
- Main method `execute(Runnable r)`
- Controls for:
 - The kind of task queue (any Channel)
 - Maximum number of threads
 - Minimum number of threads
 - "Warm" versus on-demand threads
 - Keep-alive interval until idle threads die
 - to be later replaced by new ones if necessary
 - Saturation policy
 - block, drop, producer-runs, etc

Multiple Reactor Threads

- Using Reactor Pools
 - Use to match CPU and IO rates
 - Static or dynamic construction
 - Each with own Selector, Thread, dispatch loop
- Main acceptor distributes to other reactors

```
Selector[] selectors; // also create threads
int next = 0;
class Acceptor { // ...
    public synchronized void run() { ...
        Socket connection = serverSocket.accept();
        if (connection != null)
            new Handler(selectors[next], connection);
        if (++next == selectors.length) next = 0;
    }
}
```

Using Multiple Reactors



Why JNI?

- * 使用现有的成熟的代码
- * 访问系统特性或设备
- * 代码运行速度很关键时

- * 为函数保留一个代码位置

```
public class HelloNative {  
    public native static void greeting();  
    public static void main(String[] args) {  
        greeting();  
    }  
}
```

- * 编译该Java程序HelloNative.java
- * 使用javah HelloNative来产生C语言头文件

* 编写C程序

* 编译生成DLL

vcvars32

```
cl -Ic:\jdk\include -Ic:\jdk\include\win32 -LD  
HelloNative.c -FeHelloNative.dll
```

* 在Java程序中装载DLL

```
static {
```

```
    System.loadLibrary("HelloNative");
```

```
}
```

传递基本参数

boolean	jboolean	1
byte	jbyte	1
char	jchar	2
short	jshort	2
int	jint	4
long	jlong	8
float	jfloat	4
double	jdouble	8

传递String

- * jstring
- * jstring NewStringUTF(JNIEnv*, const char[]);
- * jsize GetStringUTFLength(JNIEnv*, jstring)
- * const jbyte* GetStringUTFChars(JNIEnv*, jstring string, jboolean*)